

# Deadlock Detection

Shin Hong

# Deadlock Bugs Frequently Occur in Real World

Application	What it does	Non-Deadlock	Deadlock
MySQL	Database Server	14	9
Apache	Web Server	13	4
Mozilla	Web Browser	41	16
OpenOffice	Office Suite	6	2
Total		74	31

- In a survey on 105 real-world concurrency bugs in open-source applications, **31 out of 105 bugs are deadlock bugs** [Lu *et al.*, ASPLOS 08]


# Deadlock

- A deadlock occurs when each of a set of threads is blocked, waiting for another thread in the set to satisfy **certain condition**
  - **release shared resource**
  - **raise event**

# Resource Deadlock in Concurrent Programs

- ABBA deadlock

Thread1 () {	Thread2 () {
1: <b>lock (X)</b>	11: <b>lock (Y)</b>
2: x = ... ;	12: y = ... ;
3: <b>lock (Y)</b>	13: <b>lock (X)</b>
4: y = ... ;	14: x = ... ;
5: <b>unlock (Y)</b>	15: <b>unlock (X)</b>
6: <b>unlock (X)</b>	16: <b>unlock (Y)</b>
}	}

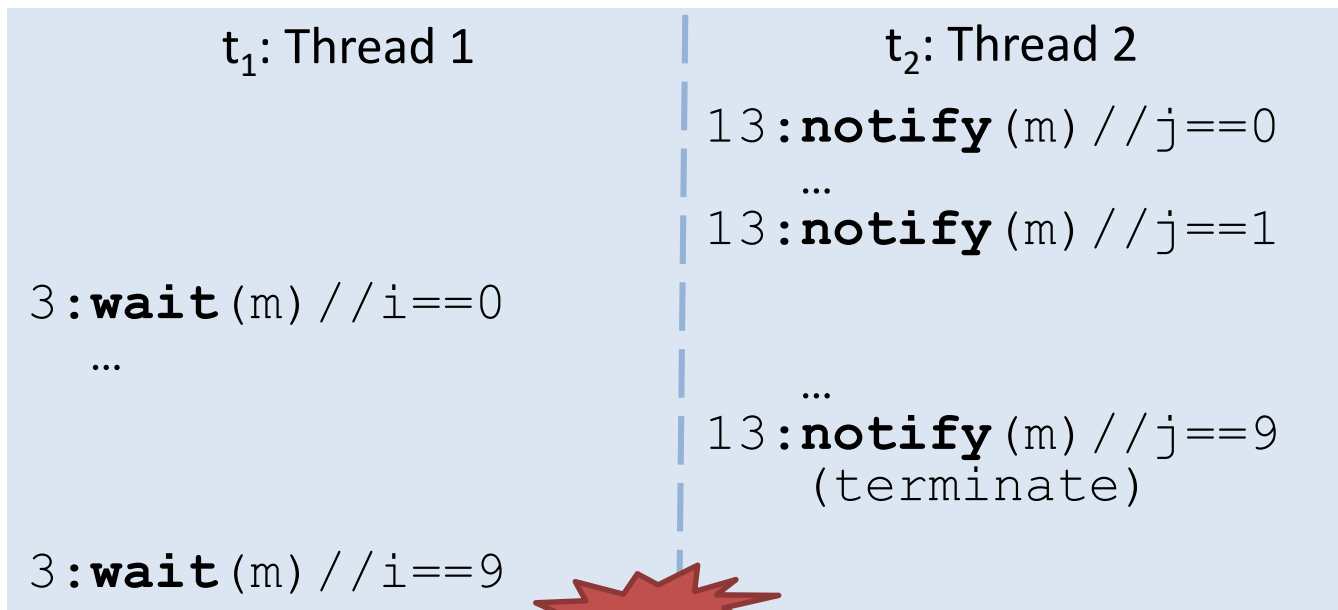


# Communication Deadlock

- Lost notify

```
Thread1 () {  
1: ...  
2: for (i=0; i<10; i++) {  
3:   wait (m) ; }  
}
```

```
Thread2 () {  
11: ...  
12: for (j=0; j<10; j++) {  
13:   notify (m) ; }  
}
```



# Finding Deadlock Bugs is Difficult

- A deadlock bug induces deadlock situations **only under certain thread schedules**
- Systems software creates a **massive number of locks** for fine-grained concurrency controls
- **Function caller-callee relation** complicates the reasoning about possible nested lockings

# Bug Detection Approach

## Resource deadlock

- Basic potential deadlock detection algorithm
- GoodLock algorithm

## Communication deadlock

- CHECKMATE: a trace program model-checking technique for deadlock detection

# Basic Potential Deadlock Detection


- Extend the cyclic deadlock monitoring algorithm
  - Cyclic deadlock monitoring algorithm (e.g. *LockDep*)
    - Monitor lock acquires and releases in runtime
    - Lock graph  $(N, E_N)$ 
      - Create a node  $n_X$  when a thread acquires lock  $X$
      - Create an edge  $(n_X, n_Y)$  when a thread acquires lock  $Y$  while holding lock  $X$
      - Remove  $n_X$ ,  $(n_X, *)$  and  $(*, n_X)$  when a thread releases  $X$
- Report deadlock when the graph has any cycle



# Cyclic Deadlock Detection Example (1/2)

```
Thread1 () {  
1: lock (X)  
2: a = ... ;  
3: lock (Y)  
4: b = ... ;  
5: unlock (Y)  
6: unlock (X)  
}  
  
Thread2 () {  
11: lock (Y)  
12: b = ... ;  
13: lock (X)  
14: a = ... ;  
15: unlock (X)  
16: unlock (Y)  
}
```

<u>t1: Thread 1</u>	<u>t2: Thread 2</u>
1: lock (X)	
2: a = ...	
	11: lock (Y)
3: lock (Y)	12: b = ...
	13: lock (X)



**Deadlock detected!**

# Cyclic Deadlock Detection Example (2/2)

```
Thread1 () {  
1: lock (X) ;  
2: a = ...  
3: lock (Y) ;  
4: b = ...  
5: unlock (Y) ;  
6: unlock (X) ;  
}  
  
Thread2 () {  
11: lock (Y) ;  
12: b = ...  
13: lock (X) ;  
14: a = ...  
15: unlock (X) ;  
16: unlock (Y) ;  
}
```

<u>t1: Thread 1</u>	<u>t2: Thread 2</u>
1: lock (X)	
2: a = ...	
3: lock (Y)	
4: b = ...	
5: unlock (Y)	
6: unlock (X)	
	11: lock (Y)
	12: b = ...
	13: lock (X)
	14: a = ...
	15: unlock (X)
	16: unlock (Y)

**No problem**

# Basic Deadlock Prediction Technique

- Potential cyclic deadlock detection algorithm [Harrow, SPIN 00]
  - Lock graph  $(N, E_N)$ 
    - Create a node  $n_x$  when a thread acquires lock  $X$
    - Create an edge  $(n_x, n_y)$  when a thread acquires lock  $Y$  while holding lock  $X$
    - ~~• Remove  $n_x, (n_x, *)$  and  $(*, n_x)$  when a thread releases  $X$~~
  - Report potential deadlocks if the resulted graph at the end of an execution has a cycle

[Harrow, SPIN 00] J. J. Harrow, Jr.: Runtime checking of multithreaded applications with Visual Threads, SPIN Workshop 2000


# Potential Cyclic Deadlock Detection Example

```

Thread1 () {
1: lock (X)
2: a = ... ;
3: lock (Y)
4: b = ... ;
5: unlock (Y)
6: unlock (X)
}

Thread2 () {
11: lock (Y)
12: b = ... ;
13: lock (X)
14: a = ... ;
15: unlock (X)
16: unlock (Y)
}
    
```

<u>t1:Thread 1</u>	<u>t2:Thread 2</u>
1: <b>lock (X)</b>	
2: a = ...	
3: <b>lock (Y)</b>	
4: b = ...	
5: <b>unlock (Y)</b>	
6: <b>unlock (X)</b>	11: <b>lock (Y)</b>
	12: b = ...
	13: <b>lock (X)</b>
	...




**Cycle → Potential deadlock**

# Basic Deadlock Prediction Technique

- The algorithm is commercialized as a SW tool VisualThreads (*HP*)
- Empirical results show that the algorithm is very effective to discover hidden deadlock bugs
- Challenge: generate many false positive

# False Positive Example#1 – Single Thread Cycle

```
Thread1() {  
1: lock(X);  
2: lock(Y);  
3: unlock(Y);  
4: unlock(X);  
  
5: lock(Y);  
6: lock(X);  
7: unlock(X);  
8: unlock(Y);}  
  
Thread2() {  
11: lock(X);  
12: unlock(X);  
  
13: lock(Y);  
14: unlock(Y);}
```



The lock graph has a cycle, but no deadlock

A cycle that consists of edges created by one thread is a false positive

## False Positive Example#2: Gate Lock

Thread1 () {

1: lock (X) ;

2: lock (Y) ;

3: lock (Z) ;

4: unlock (Z) ;

5: unlock (Y) ;

6: unlock (X) ; }

Thread2 () {

11: lock (X) ;

12: lock (Z) ;


13: lock (Y) ;

14: unlock (Y) ;

15: unlock (Z) ;

16: unlock (X) ; }

Gate lock  
(guard lock)




Cycle, but no deadlock

# False Positive Example#3: Thread Creation

```
main() {  
0:   start(f1);  
}  
  
f1() {  
1:   lock(X);  
2:   lock(Y);  
3:   unlock(Y);  
4:   unlock(X);  
5:   start(f2);  
}  
  
f2() {  
11:  lock(Y);  
12:  lock(X);  
13:  unlock(X);  
14:  unlock(Y);  
}
```

Thread segment#1

Thread segment#2



Cycle, but no deadlock



# GoodLock Algorithm[Agarwal, IBM 10]

- Extend the lock graph in the basic potential deadlock detection algorithm to consider *thread*, *gate lock*, and *thread segment*
- A cycle is *valid* (i.e., true positive) when every pair of edges  $(m_{11}, (s_{11}, t_1, G_1, s_{12}), m_{12})$ , and  $(m_{21}, (s_{21}, t_2, G_2, s_{22}), m_{22})$  in the cycle satisfies:
  - $t_1 \neq t_2$ , and
  - $G_1 \cap G_2 = \emptyset$ , and
  - $\neg(s_{12} < s_{21})$ 
    - The happens-before relation  $<$  is the transitive closure of the relation  $R$  such that  $(s_1, s_2) \in R$  if there exists the edge from  $s_1$  to  $s_2$  in the thread segment graph

[Agarwal, IBM 10] R. Agarwal et al., Detection of deadlock potential in multithreaded programs, IBM Journal of Research and Development, 54(5), 2010